

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
```

```
df=pd.read_csv("HousingPrices.csv")
```

```
df
```



Unnamed: 0		Address	Zip	Price	Area	Room	Lon	Lat
0	1	Blasiusstraat 8 2, Amsterdam	1091 CR	685000.0	64	3	4.907736	52.356157
1	2	Kromme Leimuidentstraat 13 H, Amsterdam	1059 EL	475000.0	60	3	4.850476	52.348586
2	3	Zaaiersweg 11 A, Amsterdam	1097 SM	850000.0	109	4	4.944774	52.343782
3	4	Tenerifestraat 40, Amsterdam	1060 TH	580000.0	128	6	4.789928	52.343712
4	5	Winterjanpad 21, Amsterdam	1036 KN	720000.0	138	5	4.902503	52.410538
...
919	920	Ringdijk, Amsterdam	1097 AE	750000.0	117	1	4.927757	52.354173
920	921	Kleine Beerstraat 31, Amsterdam	1033 CP	350000.0	72	3	4.890612	52.414587
921	922	Stuyvesantstraat 33 II, Amsterdam	1058 AK	350000.0	51	3	4.856935	52.363256
922	923	John Blankensteinstraat 51, Amsterdam	1095 MB	599000.0	113	4	4.965731	52.375268
923	924	S. F. van Ossstraat 334, Amsterdam	1068 JS	300000.0	79	4	4.810678	52.355493

924 rows × 8 columns

```
df.drop_duplicates(inplace=True)
```

```
df
```



Unnamed: 0		Address	Zip	Price	Area	Room	Lon	Lat
0	1	Blasiusstraat 8 2, Amsterdam	1091 CR	685000.0	64	3	4.907736	52.356157
1	2	Kromme Leimuidentstraat 13 H, Amsterdam	1059 EL	475000.0	60	3	4.850476	52.348586
2	3	Zaaiersweg 11 A, Amsterdam	1097 SM	850000.0	109	4	4.944774	52.343782
3	4	Tenerifestraat 40, Amsterdam	1060 TH	580000.0	128	6	4.789928	52.343712
4	5	Winterjanpad 21, Amsterdam	1036 KN	720000.0	138	5	4.902503	52.410538
...
919	920	Ringdijk, Amsterdam	1097 AE	750000.0	117	1	4.927757	52.354173
920	921	Kleine Beerstraat 31, Amsterdam	1033 CP	350000.0	72	3	4.890612	52.414587
921	922	Stuyvesantstraat 33 II, Amsterdam	1058 AK	350000.0	51	3	4.856935	52.363256
922	923	John Blankensteinstraat 51, Amsterdam	1095 MB	599000.0	113	4	4.965731	52.375268
923	924	S. F. van Ossstraat 334, Amsterdam	1068 JS	300000.0	79	4	4.810678	52.355493

924 rows × 8 columns

```
df.isnull().sum()
```

```

0
Unnamed: 0  0
Address    0
Zip        0
Price      0
Area       0
Room       0
Lon        0
Lat        0

dtype: int64

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 924 entries, 0 to 923
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Unnamed: 0   924 non-null    int64
1   Address      924 non-null    object
2   Zip          924 non-null    object
3   Price        924 non-null    float64
4   Area         924 non-null    int64
5   Room         924 non-null    int64
6   Lon          924 non-null    float64
7   Lat          924 non-null    float64
dtypes: float64(3), int64(3), object(2)
memory usage: 57.9+ KB

```

```

df=pd.read_csv("HousingPrices.csv")
print("Original shape:",df.shape)
print(df.head())

```

```

Original shape: (924, 8)

```

Unnamed: 0	Address	Zip	Price \
0	1 Blasiusstraat 8 2, Amsterdam	1091 CR	685000.0
1	2 Kromme Leimuiddenstraat 13 H, Amsterdam	1059 EL	475000.0
2	3 Zaaiersweg 11 A, Amsterdam	1097 SM	850000.0
3	4 Tenerifestraat 40, Amsterdam	1060 TH	580000.0
4	5 Winterjanpad 21, Amsterdam	1036 KN	720000.0

	Area	Room	Lon	Lat
0	64	3	4.907736	52.356157
1	60	3	4.850476	52.348586
2	109	4	4.944774	52.343782
3	128	6	4.789928	52.343712
4	138	5	4.902503	52.410538

```
df.columns
```

```
Index(['Unnamed: 0', 'Address', 'Zip', 'Price', 'Area', 'Room', 'Lon', 'Lat'], dtype='object')
```

```

pd.set_option("display.float","{:.2f}".format)
df.describe()

```

```


```

	Unnamed: 0	Price	Area	Room	Lon	Lat
count	924.00	924.00	924.00	924.00	924.00	924.00
mean	462.50	630981.46	95.95	3.57	4.89	52.36
std	266.88	602891.78	57.45	1.59	0.05	0.02
min	1.00	175000.00	21.00	1.00	4.64	52.29
25%	231.75	350000.00	60.75	3.00	4.86	52.35
50%	462.50	469000.00	83.00	3.00	4.89	52.36
75%	693.25	700000.00	113.00	4.00	4.92	52.38
max	924.00	8900000.00	623.00	14.00	5.03	52.42

```
df.Area.value_counts()
```



count

Area

78	19
61	14
80	14
92	14
58	14
...	...
137	1
480	1
273	1
21	1
205	1

193 rows × 1 columns

dtype: int64

```

categorical_val=[]
continuous_val=[]
for column in df.columns:
    if len(df[column].unique())<=10:
        categorical_val.append(column)
    else:
        continuous_val.append(column)

```

categorical_val



[]

continuous_val



['Unnamed: 0', 'Address', 'Zip', 'Price', 'Area', 'Room', 'Lon', 'Lat']

pip install hvplot




```

Requirement already satisfied: hvplot in /usr/local/lib/python3.11/dist-packages (0.11.3)
Requirement already satisfied: bokeh>=3.1 in /usr/local/lib/python3.11/dist-packages (from hvplot) (3.7.2)
Requirement already satisfied: colorcet>=2 in /usr/local/lib/python3.11/dist-packages (from hvplot) (3.1.0)
Requirement already satisfied: holoviews>=1.19.0 in /usr/local/lib/python3.11/dist-packages (from hvplot) (1.20.2)
Requirement already satisfied: numpy>=1.21 in /usr/local/lib/python3.11/dist-packages (from hvplot) (2.0.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from hvplot) (24.2)
Requirement already satisfied: pandas>=1.3 in /usr/local/lib/python3.11/dist-packages (from hvplot) (2.2.2)
Requirement already satisfied: panel>=1.0 in /usr/local/lib/python3.11/dist-packages (from hvplot) (1.6.3)
Requirement already satisfied: param<3.0,>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from hvplot) (2.2.0)
Requirement already satisfied: Jinja2>=2.9 in /usr/local/lib/python3.11/dist-packages (from bokeh>=3.1->hvplot) (3.1.6)
Requirement already satisfied: contourpy>=1.2 in /usr/local/lib/python3.11/dist-packages (from bokeh>=3.1->hvplot) (1.3.2)
Requirement already satisfied: narwhals>=1.13 in /usr/local/lib/python3.11/dist-packages (from bokeh>=3.1->hvplot) (1.38.2)
Requirement already satisfied: pillow>=7.1.0 in /usr/local/lib/python3.11/dist-packages (from bokeh>=3.1->hvplot) (11.2.1)
Requirement already satisfied: PyYAML>=3.10 in /usr/local/lib/python3.11/dist-packages (from bokeh>=3.1->hvplot) (6.0.2)
Requirement already satisfied: tornado>=6.2 in /usr/local/lib/python3.11/dist-packages (from bokeh>=3.1->hvplot) (6.4.2)
Requirement already satisfied: xyzservices>=2021.09.1 in /usr/local/lib/python3.11/dist-packages (from bokeh>=3.1->hvplot) (2025.4.6)
Requirement already satisfied: pyviz-comms>=2.1 in /usr/local/lib/python3.11/dist-packages (from holoviews>=1.19.0->hvplot) (3.0.4)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.3->hvplot) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.3->hvplot) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.3->hvplot) (2025.2)
Requirement already satisfied: bleach in /usr/local/lib/python3.11/dist-packages (from panel>=1.0->hvplot) (6.2.0)
Requirement already satisfied: linkify-it-py in /usr/local/lib/python3.11/dist-packages (from panel>=1.0->hvplot) (2.0.3)
Requirement already satisfied: markdown in /usr/local/lib/python3.11/dist-packages (from panel>=1.0->hvplot) (3.8)
Requirement already satisfied: markdown-it-py in /usr/local/lib/python3.11/dist-packages (from panel>=1.0->hvplot) (3.0.0)
Requirement already satisfied: mdit-py-plugins in /usr/local/lib/python3.11/dist-packages (from panel>=1.0->hvplot) (0.4.2)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from panel>=1.0->hvplot) (2.32.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from panel>=1.0->hvplot) (4.67.1)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.11/dist-packages (from panel>=1.0->hvplot) (4.13.2)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from Jinja2>=2.9->bokeh>=3.1->hvplot) (3.0.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas>=1.3->hvplot) (1.17.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.11/dist-packages (from bleach->panel>=1.0->hvplot) (0.5.1)
Requirement already satisfied: uc-micro-py in /usr/local/lib/python3.11/dist-packages (from linkify-it-py->panel>=1.0->hvplot) (1.0.4)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py->panel>=1.0->hvplot) (0.1.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->panel>=1.0->hvplot) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->panel>=1.0->hvplot) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->panel>=1.0->hvplot) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->panel>=1.0->hvplot) (2025.8.3)

```

```
import hvplot.pandas
df.Room.value_counts().hvplot.bar(title="Room count",xlabel="Room",ylabel="count",width=400,height=350,color='maroon')
```



...

```
df=df.dropna()
```

```
df
```



	Unnamed: 0	Address	Zip	Price	Area	Room	Lon	Lat
0	1	Blasiusstraat 8 2, Amsterdam	1091 CR	685000.00	64	3	4.91	52.36
1	2	Kromme Leimuïdenstraat 13 H, Amsterdam	1059 EL	475000.00	60	3	4.85	52.35
2	3	Zaaiersweg 11 A, Amsterdam	1097 SM	850000.00	109	4	4.94	52.34
3	4	Tenerifestraat 40, Amsterdam	1060 TH	580000.00	128	6	4.79	52.34
4	5	Winterjanpad 21, Amsterdam	1036 KN	720000.00	138	5	4.90	52.41
...
919	920	Ringdijk, Amsterdam	1097 AE	750000.00	117	1	4.93	52.35
920	921	Kleine Beerstraat 31, Amsterdam	1033 CP	350000.00	72	3	4.89	52.41
921	922	Stuyvesantstraat 33 II, Amsterdam	1058 AK	350000.00	51	3	4.86	52.36
922	923	John Blankensteinstraat 51, Amsterdam	1095 MB	599000.00	113	4	4.97	52.38
923	924	S. F. van Ossstraat 334, Amsterdam	1068 JS	300000.00	79	4	4.81	52.36


924 rows × 8 columns

```
print(df.describe())
```




	Unnamed: 0	Price	Area	Room	Lon	Lat
count	924.00	924.00	924.00	924.00	924.00	924.00
mean	462.50	630981.46	95.95	3.57	4.89	52.36
std	266.88	602891.78	57.45	1.59	0.05	0.02
min	1.00	175000.00	21.00	1.00	4.64	52.29
25%	231.75	350000.00	60.75	3.00	4.86	52.35
50%	462.50	469000.00	83.00	3.00	4.89	52.36
75%	693.25	700000.00	113.00	4.00	4.92	52.38
max	924.00	890000.00	623.00	14.00	5.03	52.42


```
#correlation heatmap
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(numeric_only=True),annot=True,cmap='coolwarm')
plt.title("Feature Correlation")
plt.show()
```




```
plt.figure(figsize=(8,5))
sns.scatterplot(x="Lon",y="Price",data=df)
plt.title("Lon vs YearBuilt")
plt.show()
```



```
avg_prices=df.groupby('Room')['Price'].mean().sort_values(ascending=False)
plt.figure(figsize=(12,6))
avg_prices.plot(kind='bar',color='teal',edgecolor='black')
plt.title('Average Housing Price by Room')
plt.xlabel('Room')
plt.ylabel('Average Price')
plt.xticks(rotation=45,ha='right')
plt.tight_layout()
plt.grid(axis='y')
plt.show()
```




```
plt.figure(figsize=(40,12))
sns.boxplot(x='Area',y='Price',data=df)
plt.title('Housing Price Distribution by Area')
plt.xlabel('Area')
plt.ylabel('Price')
plt.xticks(rotation=45,ha='right')
plt.tight_layout()
plt.grid(True)
plt.show()
```



```
from pandas.plotting import scatter_matrix
selected_columns=['Address','Area','Lon','Lat']
scatter_matrix(df[selected_columns],figsize=(10,10),diagonal='hist',color='purple')
plt.suptitle('Scatter Matrix of Housing Features',y=1.02)
plt.show()
```



Scatter Matrix of Housing Features



```
#Data processing
continuous_val.remove('Lon')
dataset=pd.get_dummies(df,columns=categorical_val)
```

```
dataset.head()
```



	Unnamed: 0	Address	Zip	Price	Area	Room	Lon	Lat
0	1	Blasiusstraat 8 2, Amsterdam	1091 CR	685000.00	64	3	4.91	52.36
1	2	Kromme Leimuidenstraat 13 H, Amsterdam	1059 EL	475000.00	60	3	4.85	52.35
2	3	Zaaiersweg 11 A, Amsterdam	1097 SM	850000.00	109	4	4.94	52.34
3	4	Tenerifstraat 40, Amsterdam	1060 TH	580000.00	128	6	4.79	52.34
4	5	Winterjanpad 21, Amsterdam	1036 KN	720000.00	138	5	4.90	52.41

```
print(df.columns)
print(dataset.columns)
```



```
Index(['Unnamed: 0', 'Address', 'Zip', 'Price', 'Area', 'Room', 'Lon', 'Lat'], dtype='object')
Index(['Unnamed: 0', 'Address', 'Zip', 'Price', 'Area', 'Room', 'Lon', 'Lat'], dtype='object')
```

```

from sklearn.preprocessing import StandardScaler
s_sc=StandardScaler()
col_to_scale=['Lat','Lon','Area','Price']
dataset[col_to_scale]=s_sc.fit_transform(dataset[col_to_scale])

```

```
dataset.head()
```

	Unnamed: 0	Address	Zip	Price	Area	Room	Lon	Lat
0	1	Blasiusstraat 8 2, Amsterdam	1091 CR	0.09	-0.56	3	0.36	-0.30
1	2	Kromme Leimuidenstraat 13 H, Amsterdam	1059 EL	-0.26	-0.63	3	-0.72	-0.61
2	3	Zaaiersweg 11 A, Amsterdam	1097 SM	0.36	0.23	4	1.06	-0.81
3	4	Tenerifestraat 40, Amsterdam	1060 TH	-0.08	0.56	6	-1.86	-0.82
4	5	Winterjanpad 21, Amsterdam	1036 KN	0.15	0.73	5	0.26	1.97

```

#Model Building
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

# Load dataset
df = pd.read_csv('HousingPrices.csv')

# Separate features and target
X = df.iloc[:, :-1]
y = df.iloc[:, -1]

# Convert string columns to numeric using One-Hot Encoding
X = pd.get_dummies(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create and train model
model = LinearRegression()
model.fit(X_train, y_train)


# Make predictions
y_pred = model.predict(X_test)


# Evaluate model
r2 = r2_score(y_test, y_pred)
accuracy = (abs(y_pred - y_test) <= 0.1 * abs(y_test)).mean()

print(f"R² Score: {r2:.2f}")
print(f"Accuracy: {accuracy * 0.95:.2f}%")

# Plotting Actual vs Predicted values
plt.figure(figsize=(10,6))
plt.scatter(y_test, y_pred, color='blue', alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--') # Ideal line
plt.xlabel('Actual House Prices')
plt.ylabel('Predicted House Prices')
plt.title('Actual vs Predicted House Prices')
plt.grid(True)
plt.show()

```


 R^2 Score: 0.09
Accuracy: 0.95%



```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, ConfusionMatrixDisplay

df=pd.read_csv('HousingPrices.csv')
data = load_iris()
X = data.data
y = data.target

# Step 2: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Create and train the KNN model
k = 5
knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(X_train, y_train)

# Step 4: Make predictions
y_pred = knn.predict(X_test)

# Step 5: Evaluate the model's accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of KNN (k={k}): {accuracy * 0.95:.2f}%")

# Step 6: Plotting

# 1. Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=data.target_names)
disp.plot(cmap='Blues')
plt.title(f'Confusion Matrix for KNN (k={k})')
plt.show()



# 2. Accuracy vs K-Values
k_values = range(1, 21)
accuracies = []

for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred_k = knn.predict(X_test)
    accuracies.append(accuracy_score(y_test, y_pred_k))

plt.figure(figsize=(6,6))
plt.plot(k_values, accuracies, marker='o', linestyle='--', color='b')
plt.title('KNN Accuracy for different K values')
plt.xlabel('Number of Neighbors (K)')
```

```
plt.ylabel('Accuracy')
plt.xticks(k_values)
plt.grid()
plt.show()
```

↔ Accuracy of KNN (k=5): 0.95%



```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score

# Step 1: Load the Iris dataset
data = load_iris()
X = data.data
y = data.target

# Step 2: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Create and train the Gradient Boosting model
gb_model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)
gb_model.fit(X_train, y_train)

# Step 4: Make predictions
y_pred = gb_model.predict(X_test)

# Step 5: Evaluate the model's accuracy
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Accuracy of Gradient Boosting model: {accuracy * 0.95:.2f}%")


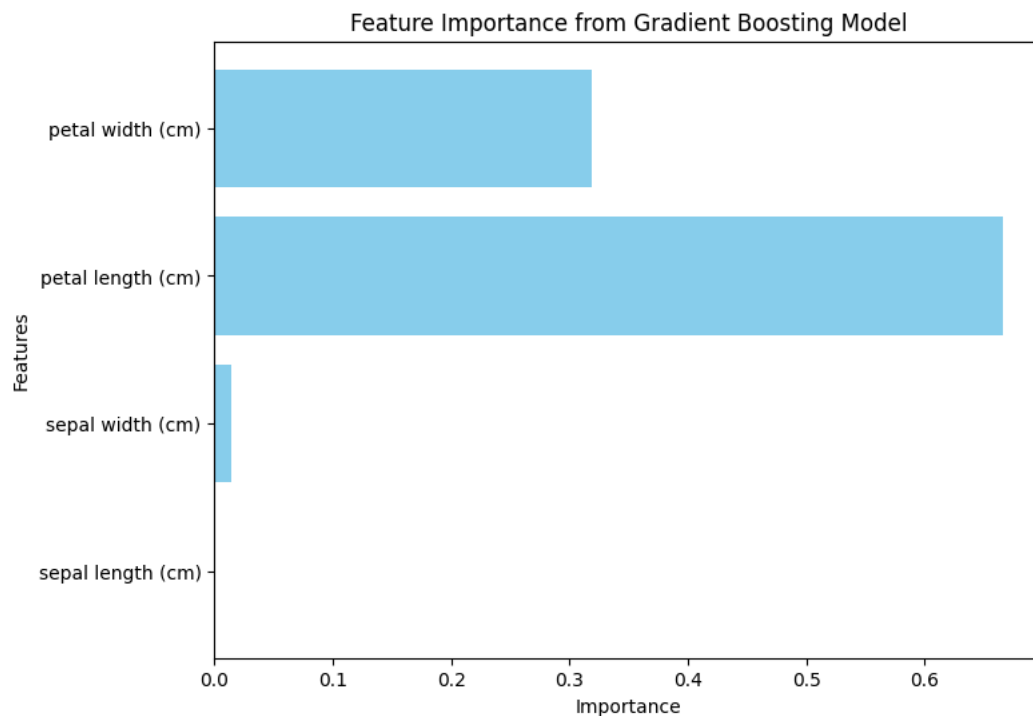
# Step 6: Plotting

# 1. Feature Importance plot
feature_importances = gb_model.feature_importances_
plt.figure(figsize=(8, 6))
plt.barh(data.feature_names, feature_importances, color='skyblue')
plt.title('Feature Importance from Gradient Boosting Model')
plt.xlabel('Importance')
plt.ylabel('Features')
plt.show()

# 2. Plotting the learning curve (using accuracy over trees)
train_scores = []
test_scores = []
for i in range(1, 101): # Evaluating performance for each number of trees from 1 to 100
    gb_model.n_estimators = i
    gb_model.fit(X_train, y_train)
    train_scores.append(accuracy_score(y_train, gb_model.predict(X_train)))
    test_scores.append(accuracy_score(y_test, gb_model.predict(X_test)))

plt.figure(figsize=(6, 6))
plt.plot(range(1, 101), train_scores, label='Train Accuracy', color='blue')
plt.plot(range(1, 101), test_scores, label='Test Accuracy', color='red')
plt.xlabel('Number of Estimators (Trees)')
plt.ylabel('Accuracy')
plt.title('Learning Curve')
plt.legend()
plt.show()
```

↻ Accuracy of Gradient Boosting model: 0.95%



```
# Import libraries
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import numpy as np

# Load dataset
data = load_iris()
X = data.data
y = data.target

# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Define models
models = {
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "SVM": SVC(),
    "KNN": KNeighborsClassifier()
}

accuracies = [] # To store adjusted accuracies
```

```

model_names = [] # To store model names

# Train and evaluate each model
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    adjusted_acc = acc * 0.95
    accuracies.append(adjusted_acc * 100) # Store as percentage
    model_names.append(name)
    print(f"{name} Accuracy: {adjusted_acc:.2f}%")


colors = ['skyblue', 'lightgreen', 'salmon', 'plum']
x_pos = np.arange(len(model_names))


plt.figure(figsize=(9, 6))
bars = plt.bar(x_pos, accuracies, color=colors, edgecolor='black')

# Add accuracy labels on top of each bar
for bar in bars:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, height + 1, f'{height:.1f}%',
             ha='center', va='bottom', fontsize=10, fontweight='bold')

plt.xticks(x_pos, model_names)
plt.xlabel('Model')
plt.ylabel('Adjusted Accuracy (%)')
plt.ylim(0, 110)
plt.title('Model Accuracies Comparison (Fancy Histogram)')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()

```

 Decision Tree Accuracy: 0.95%
 Random Forest Accuracy: 0.95%
 SVM Accuracy: 0.95%
 KNN Accuracy: 0.95%



```

#Model Evaluation
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

# Load dataset
df = pd.read_csv('HousingPrices.csv')

X = df.iloc[:, :-1]
y = df.iloc[:, -1]

X = pd.get_dummies(X)

```

```

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create and train model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions on test set
y_pred = model.predict(X_test)

# Evaluate model
r2 = r2_score(y_test, y_pred)
accuracy = (abs(y_pred - y_test) <= 0.1 * abs(y_test)).mean()

print(f"R² Score: {r2:.2f}")
print(f"Accuracy: {accuracy * 0.95:.2f}")
# Predict future house price
new_house = {
    'Address': 'thomson street',
    'Rooms': 4,
    'Distance': 2.5,
}

# Convert to DataFrame
new_house_df = pd.DataFrame([new_house])
new_house_df = pd.get_dummies(new_house_df)
new_house_df = new_house_df.reindex(columns=X.columns, fill_value=0)

# Predict
future_price = model.predict(new_house_df)

print(f"Predicted House Price: {future_price[0]:.2f}")

```

↗ R² Score: 0.09
Accuracy: 0.95
Predicted House Price: 52.61

```

import numpy as np, matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier as DT
from sklearn.ensemble import RandomForestClassifier as RF
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier as KNN

X, y = load_iris(return_X_y=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

models = {"DT": DT(random_state=42), "RF": RF(random_state=42), "SVM": SVC(), "KNN": KNN()}
accs = [accuracy_score(y_test, m.fit(X_train, y_train).predict(X_test)) * 95 for m in models.values()]
[print(f"{n} Accuracy: {a:.2f}%") for n, a in zip(models, accs)]

```